

Recognizing Sleep Stages with Wearable Sensors in Everyday Settings

Ulrich Reimer¹, Sandro Emmenegger¹, Edith Maier¹, Zhongxing Zhang²
and Ramin Khatami²

¹*Institute for Information and Process Management, University of Applied Sciences St. Gallen, Switzerland*

²*Center for Sleep Research and Sleep Medicine, Clinic Barmelweid, Switzerland*

Keywords: Sensor Data Analysis, Deep Learning, Sleep Stage Recognition, Behavioural Change Support.

Abstract: The paper presents results from the SmartSleep project which aims at developing a smartphone app that gives users individual advice on how to change their behaviour to improve their sleep. The advice is generated by identifying correlations between behaviour during the day and sleep architecture. To this end, the project addresses two sub-tasks: detecting a user's daytime behaviour and recognising sleep stages in an everyday setting. The focus of the paper is on the second task. Various sensor devices from the consumer market were used in addition to the usual PSG sensors in a sleep lab. An expert assigned a sleep stage for every 30 seconds. Subsequently, a sleep stage classifier was learned from the resulting sensor data streams segmented into labelled sleep stages of 30 seconds each. Apart from handcrafted features we also experimented with unsupervised feature learning based on the deep learning paradigm. Our best results for correctly classified sleep stages are in the range of 90 to 91% for Wake, REM and N3, while the best recognition rate for N2 is 83%. The classification results for N1 turned out to be much worse, N1 being mostly confused with N2.

1 INTRODUCTION

Sleep quality is associated with health, wellbeing and quality of life. Sleep disorders, however, are widespread and often coincide with chronic health problems such as diabetes, hypertension, obesity as well as cardiovascular and psychiatric diseases such as depression. According to a recent survey (Tinguely et al., 2014), about 20% of people in Switzerland suffer from sleep disorders. About 28% of those affected were taking sleeping pills on a regular basis. Approximately 80% of patients with depression also complained about sleep disorders which can be considered predictors of future depression. According to a meta analysis of over 20 published longitudinal studies between 1980 and 2010, insomnia doubles the risk of suffering from depression (Baglioni et al., 2011).

Overall, experts agree that the prevalence of sleep disorders such as obstructive sleep apnoea (OAS) or daytime sleepiness tends to be underestimated. Sleep disorders therefore often remain undiagnosed and untreated even though they are a significant cause of morbidity and mortality (Hossain and Shapiro, 2002); for a detailed review of sleep disorders sleepiness, see (Panossian and Avidan, 2009).

Traditionally, sleep disorders are investigated in sleep laboratories by means of polysomnography (PSG) as well as by actigraphic assessment. A polysomnogram or sleep study usually involves the measurement of brain activity through the electroencephalogram (EEG), muscular activity (EMG) and eye movements (EOG). Other parameters monitored include oxygen saturation, respiratory effort, cardiac activity as well as sound and movement activity.

But not only is such sleep monitoring costly, it also removes people from their normal sleeping environment and prevents repeated or longitudinal studies. Increasingly, home sleep recording systems are coming on the market which aim to reduce the financial cost and reach a larger population (Subramanian et al., 2011). However, without medical or technical training, people often fail to place the sensors in the correct positions which results in inconclusive data. Even if done correctly, the challenge remains of actually analysing or scoring the data, which requires specific expertise.

More recently, smart watches, fitness trackers as well as sensors built into a smartphone offer new opportunities for continuous monitoring in every-day settings. Sensors and wearables can capture data

about people's rest and activity patterns. Most devices use accelerometers for tracking movements during the night from which they derive information on sleep architecture and sleep quality. Some devices take additional vital parameters into account, such as heart rate and skin conductance.

However, the tracking devices and sleep screening apps currently available cannot compete with the accuracy of clinical sleep laboratories. At best, they are able to distinguish between waking time and sleep time. When we compared different devices that claim to distinguish sleep phases we found little match between the identified sleep stages. Besides, according to a review of current sleep screening applications conducted by Behar and his team (Behar et al., 2013) none of the existing sleep monitoring applications available for smartphones with the exception of simple questionnaires, is based on scientific evidence.

The ability to reliably detect sleep stages and thus monitor sleep architecture on a daily basis in a home setting is a *prerequisite for*

- finding individual correlations between behaviour during the day and sleep architecture,
- measuring the effects interventions and behavioural changes have on sleep architecture, e.g. for monitoring therapeutical effects of daytime activity of patients suffering from depression, or the effects of interventions aimed at stabilising sleep/wake phases especially with older people,
- monitoring the effects of individual cognitive behavioural therapy for insomnia, especially the elements of sleep restriction and stimulus control,
- enabling the recognition and quantification of the effects of activity and movement as well as sleep quality on the rehabilitation process.

Given the shortcomings of existing solutions for monitoring sleep architecture in a home setting, the SmartSleep¹ Project set out to achieve the following *objectives* (cf. Fig.1):

- use data captured with wearable sensors to identify and record sleep stages with an accuracy approximating a clinical polysomnography,
- based on the above, develop a low-cost monitoring solution for capturing sleep architecture at home over a longer period of time,

¹The SmartSleep project is funded by the International Bodensee Hochschule. The consortium includes the Universities of Applied Sciences of St. Gallen, of Vorarlberg and of Constance, the Center for Sleep Research and Sleep Medicine at the Swiss Clinic Barmelweid and the two SMEs Biovation and myVitali.

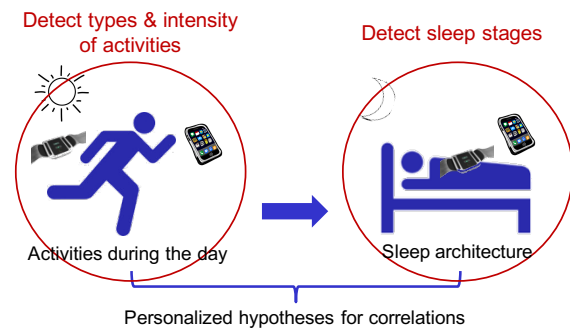


Figure 1: The objectives of the SmartSleep project.

- collect data about a user's daytime behaviour and environmental factors to identify possible correlations with sleep architecture as well as a person's perceived quality of sleep.

The last objective is the ultimate goal of the project, which is expected to open up the possibility to track a person's response to simple behavioural interventions e.g. more physical activity or more exposure to ambient light during the day, and based on these insights to automatically generate individual advice for behavioural changes.

2 APPROACH

The main project objective is to develop a smartphone app that gives users individual advice on how to change their behaviour to improve their sleep. The advice is generated by identifying significant correlations between behaviour during the day and sleep architecture. Additionally, a subjective assessment of sleep quality can be done using a questionnaire. Such correlations provide useful indications about which kinds of behaviour have a positive or a negative effect on sleep quality.

The impact of behaviour changes can subsequently be measured by repeating the measurements of sleep architecture and a subjective assessment of sleep quality.

The correlations discovered by data mining are highly personal because individuals differ greatly with regard to what may promote and what may hinder sleep. Whereas for one person a walk in the evening is very conducive to a good night's sleep, someone else may get too agitated. This is why in the project we focus on recognising patterns that apply to a specific individual rather than on statistical correlations in a population.

To achieve the project goals we address two sub-tasks. Firstly, we are developing a component for activity recognition in order to detect a user's daytime

behaviour (Sec.2.1). Secondly, we are developing a component for recognising sleep architecture, i.e. the sequence and frequency of the various types of sleep stages in the course of a night (Sec.2.2).

2.1 Activity Detection

At present, we are considering the following features for characterizing daytime behaviour:

- **Elementary activities** such as walking, running, cycling: Elementary activities are detected from the data of **one or two accelerometers** a user is wearing. A detection algorithm with an accuracy of more than 90% has been developed based on algorithms published in the literature – see e.g. (Kwapisz et al., 2011; Alsheikh et al., 2015).
- **Complex activities** such as household chores, working in the garden or kitchen: In the project we are currently developing a classifier for detecting complex activities using data mining techniques (Sohm, 2016). The classifier can be trained by users individually by giving feedback on their activities. In this way, activity detection is tailored to those activities relevant for each individual user. The classifier makes use of accelerometer data taken from the smartphone or from accelerometers worn at the wrist and/or ankle.
- **Body postures** (standing, sitting, lying): **One** of the accelerometers we are using has built-in posture detection.
- **Stress level**: Stress level is captured by the app developed in our SmartCoping project (Reimer et al., 2017)
- **Activity index** associated with times of the day (morning, afternoon, evening): The activity index is calculated from the duration and intensity of movement during a given time period.

Further data sources for detecting daytime behaviour might be added in the future, e.g. data from smart metering of electricity and water consumption, which would allow a quite detailed monitoring of a person's behaviour at home. Of course, privacy issues need to be taken special care of.

2.2 Recognition of Sleep Stages from Wearable Sensors

The automatic detection of sleep stages from sensor data is a goal that many researchers are currently pursuing. Most existing approaches work on the polysomnography data generated in a sleep lab, i.e. EEG,

EOG and EMG (see e.g. (Långkvist et al., 2012; Herrera et al., 2013; Shi et al., 2015)). So far only very few papers have been published on detecting sleep stages from wearable sensors developed for the consumer market. Some of them use other consumer sensors as the ground truth rather than a clinical gold standard such as PSG against which they evaluate their systems and algorithms (Gu et al., 2014; Rahman et al., 2015).

Among the few papers that have reported the use of sensors suited to a home setting, only a small fraction has actually validated their results against a clinical gold standard such as **PSG or the Rechtschaffen and Kales method (R-K method)**. **Automatic sleep stage recognition based on heart rate and body movement was investigated in** (Kurihara and Watanabe, 2012) **and the accuracy of their system compared with the R-K method**. (O'Hare et al., 2015) discuss the detection of sleep and waking time by various motion sensors and compare them with PSG measurements taken in parallel. They were not concerned with the detection of different sleep stages, however.

For our SmartSleep project we could not use the sleep stage recognition services of any of the existing body sensors that offer sleep stage detection because the sleep stage recognition implemented has been shown to be rather unreliable. Many of them even have difficulty in distinguishing sleep from waking phases with sufficient accuracy (Kolla et al., 2016). These findings coincide with the results of our own experiments which compared two of those devices and found nearly no correlation between detected sleep stages at all. Therefore it was necessary to develop our own recognition algorithm.

Recognition algorithms are not handcrafted but obtained by learning a classifier for each sleep stage. **For this we need a) appropriate input data from which to learn the sleep stage classifiers, b) an appropriate learning algorithm**. The input data was provided by the **clinical project partner** who has taken **measurements with our consumer sensors in parallel to classical PSG**. Using data mining algorithms we then were able to correlate the sleep stages as recorded in the PSG hypnogram with specific patterns in the data from the consumer sensors. The patterns so identified can then be used to segment sensor data streams into sleep stages. The following section describes our approach in more detail, first explaining the experimental set-up (Sec.3.1) and then distinguishing between using *handcrafted features* (Sec.3.2) and features learned via a *deep belief network* (Sec.3.3).

3 LEARNING SLEEP STAGE CLASSIFIERS

3.1 Experimental Set-up

We have been experimenting with several kinds of wearable sensors and finally focused on the following two:

- Zephyr BioHarness 3² chest strap with a reporting frequency of 1 Hz for the channels: heart rate, breathing rate, breathing rate amplitude, ECG amplitude as well as minimum and peak levels of the vertical, lateral and sagittal axes
- two MSR 145B³ accelerometers – one at the wrist and one at the ankle, with a sampling frequency of 51.2 Hz

These sensors were given to 26 healthy volunteers in addition to the usual PSG sensors in the sleep lab of the clinical project partner. For each person the sleep stages ('Wake', 'REM', 'N1', 'N2', 'N3')⁴ were labelled by experts according to the gold standard of the AASM classification⁵. This resulted in sensor data streams segmented into labelled sleep stages of 30 seconds each, from which the sleep stage classifiers were subsequently learned. We used a Random Forest classifier, an ensemble learning approach which has also been used by other researchers for learning sleep stage classifiers and has shown to be superior to an SVM ensemble by (Radha et al., 2014).

The sleep stage classifiers learned from the data of the 26 healthy volunteers will from now on serve as a baseline. Further classifiers will be learned for patients with specific diagnoses.

We used the Weka libraries⁶ to learn the Random Forest classifier. The data processing pipeline was implemented with the Matlab language in an object-oriented architecture. The classes and processing stages were inspired by the pipes and filter patterns described in (Buschmann et al., 2013). This enables us to set up new experiments in a fast and flexible way by appropriately combining data file readers, interpolation and data merging stages, filtering, feature construction and classification steps. The execution environment is Matlab 2016b.

²www.zephyranywhere.com

³www.msr.ch

⁴REM corresponds to rapid eye movement sleep, while N1 to N3 correspond to progressively deeper stages of sleep, N1 standing for light sleep, N3 for deep sleep.

⁵www.esst.org/adds/ICSD.pdf

⁶weka.wikispaces.com

The quality of the learned classifiers not only depends on the chosen algorithm and its parameter settings but above all depends on the features being used. Especially in the case of learning from sensor data, identifying significant features is a critical and difficult task. We have experimented with handcrafted features (see Fig. 2 and Sec.3.2) as well as with unsupervised feature learning based on the deep learning paradigm (see Fig. 2 and Sec.3.3).

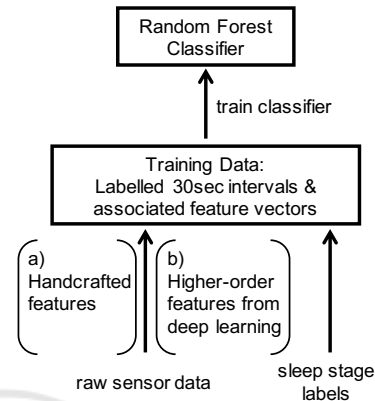


Figure 2: Feature generation and classification pipeline.

3.2 Handcrafted Features

Finding significant features usually involves much experimentation, in particular in the case of sensor data streams. We did a literature review to identify features that worked for other researchers. We looked into handcrafted features used for recognising sleep stages, e.g. (Panagiotou et al., 2015), as well as for recognising activity types, e.g. (Alsheikh et al., 2015). Based on the literature review, we decided to use the following functions to calculate the features from the sensor raw data:

- energy (sum of power at each frequency),
- max frequency,
- root mean square of sensor channel values,
- skewness (asymmetry of the probability distribution relative to its mean),
- standard deviation,
- vector norm (length of vector of sensor channel values).

More systematic experimentation with other features will be needed including an analysis of the impact of each feature on the learned classifiers. Based on the insights gained we might derive a more appropriate feature set resulting in classifiers that are more accurate.

The handcrafted features are functions which aggregate the raw data of each 30 second window and each sensor channel into a value. For the ten channels of the Zephyr sensor and the six channels of the two MSR accelerometers, this results in a feature vector of 96 components per 30 second sleep stage event. With these features a Random Forest classifier of 99 trees was learned. The confusion matrices in Tables 1, 2 and 3 show the classification accuracy of the learned classifiers based on a *tenfold cross-validation*.

Table 1: Confusion matrix: Handcrafted features MSR.

Instances:	20271				
Correctly Classified:	15434	76.1%			
		Predicted			
%	REM	Wake	N1	N2	N3
REM	83.2	3.9	9.0	3.6	0
Wake	0.9	84.6	10.4	3.7	0
N1	7.4	14.7	49.3	27.6	1.0
N2	1.5	4.0	13.3	76.9	4.4
N3	0	2.6	1.6	7.3	88.2

Table 2: Confusion matrix: Handcrafted features Zephyr.

Instances:	4891				
Correctly Classified:	3797	77.6%			
		Predicted			
%	REM	Wake	N1	N2	N3
REM	87.2	1.9	2.7	6.6	1.7
Wake	0.8	86.7	5.4	6.2	1.0
N1	12.6	16.4	26.6	40.7	3.7
N2	2.7	3.4	7.5	79.0	7.3
N3	0.8	1.9	0.6	6.1	90.7

Table 3: Confusion matrix: Handcrafted features MSR & Zephyr.

Instances:	4695				
Correctly Classified:	3790	80.7%			
		Predicted			
%	REM	Wake	N1	N2	N3
REM	91.0	1.4	2.1	4.2	1.2
Wake	0	90.3	4.5	4.1	0.8
N1	8.1	16.3	32.9	41.0	1.7
N2	1.9	2.0	7.6	83.1	5.4
N3	0	2.2	0	5.7	91.3

3.3 Unsupervised Feature Learning Using a Deep Belief Network

Feature engineering is a labour-intensive task. Inspired by the recent enthusiasm about deep learning (Bengio et al., 2012) we decided to find out how learning a Random Forest classifier using handcrafted features related to one using features learned via *deep learning*. Especially in the context of learning sleep stage classifiers, (Långkvist et al., 2012) have already




Figure 3: Structure of a deep belief network of two stacked Restricted Boltzmann Machines.

shown that unsupervised feature learning with deep learning is promising. We followed a similar approach and automatically *derived higher-order features from the raw data of the sensors by applying a deep belief network (DBN) built from stacked Restricted Boltzmann Machines (RBM)* – cf. Fig.3. These higher-order features reflect significant patterns in the underlying raw data and are therefore well suited as features for a classifier.

We used the open source Deep Belief Network Matlab implementation DeeBNet⁷ (Keyvanrad and Homayounpour, 2014). The input vector to the DBN, i.e. the raw sensor data for each 30 second sleep stage event, is constructed as follows:

- For the Zephyr chest strap we have 10 channels with a sampling frequency of 1 Hz. This amounts to 300 components in the input vector.
- For the two MSR accelerometers we have 6 channels in total with an interpolated sampling frequency of 20 Hz. This amounts to 3600 components in the input vector.

For the approximation of the log-likelihood gradient the one-step contrastive divergence (CD) method as proposed by (Hinton, 2002; Carreira-Perpinan and Hinton, 2005) was applied. As part of future fine-tunings we might experiment with other approximation methods and parameters for the RBMs in the DBN.

Configuring a DBN and finding a good topology requires both expertise and experimentation (Hinton, 2012). We experimented with various combinations of numbers of hidden layers and numbers of hidden units for each layer and also used different numbers of learning epochs. It turned out that the results in terms of classifier accuracy did not change significantly.

The accuracy of the learned Random Forest classifier based on the features learned by the DBN is

⁷ceit.aut.ac.ir/~keyvanrad/DeeBNet%20Toolbox.html

shown by the confusion matrices in Tables 4, 5 and 6. They are based on a DBN with two hidden layers, the second one being larger than the first layer. The results are discussed in the following section.

Table 4: Confusion matrix: DBN-created features MSR.

Layer 1:	886 hidden units, 50 epochs				
Layer 2:	14184 hidden units, 150 epochs				
Instances:	17803				
Correctly Classified:	13820	77.6%			
		Predicted			
%	REM	Wake	N1	N2	N3
REM	86.6	2.0	7.4	3.6	0
Wake	1.2	83.5	10.4	4.2	0.7
N1	7.9	13.8	48.5	28.0	1.8
N2	1.4	2.3	11.2	79.7	5.4
N3	0	1.6	1.5	6.7	89.8

Table 5: Confusion matrix: DBN-created features Zephyr.

Layer 1:	1320 hidden units, 150 epochs				
Layer 2:	1320 hidden units, 150 epochs				
Instances:	14760				
Correctly Classified:	9702	65.7%			
		Predicted			
%	REM	Wake	N1	N2	N3
REM	76.1	2.4	7.6	10.8	3.0
Wake	3.1	66.9	13.9	12.7	3.4
N1	11.4	9.3	36.5	37.2	5.5
N2	5.4	3.0	13.3	67.6	10.8
N3	2.1	2.3	2.7	13.1	79.8

Table 6: Confusion matrix: DBN-created features MSR & Zephyr.

Layer 1:	1965 hidden units, 150 epochs				
Layer 2:	7860 hidden units, 150 epochs				
Instances:	12103				
Correctly Classified:	8978	74.2%			
		Predicted			
%	REM	Wake	N1	N2	N3
REM	86.1	1.3	6.8	4.8	1.0
Wake	1.5	76.9	11.6	8.5	1.5
N1	10.9	9.8	41.0	35.2	3.2
N2	2.8	2.2	10.3	76.7	8.0
N3	0.8	1.7	1.6	9.8	86.2

4 DISCUSSION OF RESULTS

Our results are significantly better than those reported by other researchers who also used sensors suitable for home settings. For example, (Kurihara and Watanabe, 2012) achieved a mean of correctly classified sleep stages of 56.2% when comparing their approach against the R-K method with 5 distinct sleep stages. For distinguishing between sleep and waking

time (O'Hare et al., 2015) presented a classifier with a mean number of correct classifications of approximately 85% against PSG measurements. Our best result for identifying waking time against any of the sleep stages is 90.3% (see Table 3). (Borazio et al., 2014) used a wrist-worn accelerometer to detect sleep and wake phases and reported a precision of 79% for detecting sleep and 75% for wake phases against PSG measurements taken in parallel.

(Gu et al., 2014) presented a classifier and evaluated it against another consumer device as a reference point and achieved 63.7% of correctly classified REM stages and 60% of correctly classified N3 stages against that device. (Rahman et al., 2015) presented a classifier that was compared against two other consumer devices and achieved 80.5% when distinguishing REM vs. non-REM stages and 89.3% when distinguishing sleep from waking times.

When using handcrafted features our experimental set-up of two accelerometers achieves a surprisingly high overall recognition rate which is comparable with that of the Zephyr chest strap. When using features learned by the DBN, recognition with the chest strap deteriorates (Table 5). The recognition rate with the accelerometers, however, stays about the same.

In all cases one would have expected the chest strap, which delivers heart rate and breathing rate in addition to accelerometer data (although measured only at the chest), to be superior to accelerometers only. The reasons for the weak performance of the chest strap are twofold: First, we had many misreadings due to the electrodes losing contact whenever a person moved. We therefore had to filter out a great proportion of the sensor data stream, which reduced the number of learning examples to about a quarter of those we had available for the accelerometers. Second, the chest strap we used delivers a reading for heart rate and breathing rate only every second, although it works internally with 250 Hz for the ECG signal. For the built-in accelerometer we only get the minimum and peak values of the last second. Especially in the case of deep learning the low resolution on all the channels impedes the learning of strong features.

The fact that recognition rates with the two accelerometers are about the same when using handcrafted features and when using learned features show that deep learning works quite nicely, i.e. without any experience which kind of features work best we achieved a high recognition rate right from the start.

Besides, it is worth mentioning that in all cases the recognition rates for sleep stage N1 were particularly weak, N1 primarily being confused with N2. Whilst

a human expert can quite easily distinguish N1 from N2 it seems that the sensor data we used either do not provide appropriate criteria for reliably recognising N1 or lack sufficient resolution in the case of the chest strap.

We also looked at the recognition rates for individual persons and how much they vary. To this end, we trained classifiers without the data from preselected persons and then classified the sensor data of those individuals using the classifier. We did this for 14 persons for the combination of the MSR and Zephyr sensors. For the 14 persons four recognition rates were above 80%, six were between 70% and 80%, three between 60% and 70% and one below that. The best recognition rate over all sleep stages was 82.1% while the worst was 56.1%, which was mainly due to a very low recognition rate of 27.3% of REM events. The associated person had nearly no REM events at all (only 18 as compared to 100 to 280 REM events for the other test persons) and they might have been atypical. We will have to investigate this drop in recognition rate more closely.

5 OUTLOOK

Because of our disappointing experience with the chest strap we have been looking for alternatives. While there are other chest straps that offer higher resolutions there would still be the problem of misreadings and artifacts because movement in bed often causes a loss of electrode contact. We have therefore begun to experiment with a different sensor from our project partner Biovotion which is positioned at the upper arm. Besides skin temperature and accelerometer data it provides three opto-electronic sensors for different wavelengths which measure the absorption by the tissue. First results are encouraging because they show recognition rates of already 77% for N2, 84% for N3 and of 77% for REM even for a small amount of learning data of 4069 labeled sleep stages. We are currently collecting more data and expect to achieve higher recognition rates than with the chest strap.

Concerning the deep belief network, we have experimented with different numbers of hidden layers and different numbers of nodes per layer but recognition rates stayed about the same. While topology does not seem to have a critical influence there are many more parameters to experiment with (Hinton, 2012), which will be one of our next tasks. Further improvements of recognition rates might come from considering transition probabilities between sleep phases, which is something we are currently looking into as

well.

The sleep stage classifiers we have learned so far are for healthy individuals but we will also learn further classifiers for patients with specific diagnoses. We also plan to learn additional classifiers from the data of both healthy persons and patients put together. Our aim is to find out if it is possible to come up with classifiers that work for all people – healthy or ill – or find out to which extent they need to be specific for certain groups of people.

Finally, we are planning to integrate the results of the SmartSleep project into our framework of behavioural change support systems (Reimer and Maier, 2016) as we are convinced that giving personalised advice is a major factor in supporting for behavioural changes.

ACKNOWLEDGEMENTS

We would like to thank the anonymous reviewers of an earlier version of this paper for their helpful comments.

REFERENCES

- Alsheikh, M. A., Selim, A., Niyato, D., Doyle, L., Lin, S., and Tan, H.-P. (2015). Deep activity recognition models with triaxial accelerometers. *arXiv:1511.04664*.
- Baglioni, C., Battagliese, G., Feige, B., Spiegelhalter, K., Nissen, C., Voderholzer, U., Lombardo, C., and Riemann, D. (2011). Insomnia as a predictor of depression: A meta-analytic evaluation of longitudinal epidemiological studies. *Journal of Affective Disorders*, 135(1–3):10–19.
- Behar, J., Roebuck, A., Domingos, J. S., Geder, E., and Clifford, G. D. (2013). A review of current sleep screening applications for smartphones. *Physiological measurement*, 34(7):R29.
- Bengio, Y., Courville, A. C., and Vincent, P. (2012). Unsupervised feature learning and deep learning: A review and new perspectives. *CoRR*, abs/1206.5538.
- Borazio, M., Berlin, E., Kckyildiz, N., Scholl, P., and Laerhoven, K. V. (2014). Towards benchmarked sleep detection with wrist-worn sensing units. In *2014 IEEE Int. Conf. on Healthcare Informatics*, pages 125–134.
- Buschmann, F., Meunier, R., Rohnert, H., Sommerlad, P., and Stal, M. (2013). *Pattern-Oriented Software Architecture. A System of Patterns*. Wiley.
- Carreira-Perpinan, M. A. and Hinton, G. E. (2005). On contrastive divergence learning. *Artificial Intelligence and Statistics*, 2005.
- Gu, W., Yang, Z., Shangguan, L., Sun, W., Jin, K., and Liu, Y. (2014). Intelligent sleep stage mining service with smartphones. In *Proc. ACM Int. Joint Conf. on Pervasive and Ubiquitous Computing*, pages 649–660.

- Herrera, L. J., Fernandes, C. M., Mora, A. M., Migotina, D., Largo, R., Guillén, A., and Rosa, A. C. (2013). Combination of heterogeneous EEG feature extraction methods and stacked sequential learning for sleep stage classification. *International Journal of Neural Systems*, 23(03):1350012.
- Hinton, G. E. (2002). Training products of experts by minimizing contrastive divergence. *Neural Computation*, 14(8):1771–1800.
- Hinton, G. E. (2012). A practical guide to training Restricted Boltzmann Machines. In Montavon, G., Orr, G. B., and Müller, K.-R., editors, *Neural networks: tricks of the trade*, pages 599–619. Springer.
- Hossain, J. L. and Shapiro, C. M. (2002). The prevalence, cost implications, and management of sleep disorders: an overview. *Sleep and Breathing*, 6(02):085–102.
- Keyvanrad, M. A. and Homayounpour, M. M. (2014). *A brief survey on deep belief networks and introducing a new object oriented toolbox (DeeBNet)*. Technical Report, Laboratory for Intelligent Multimedia Processing, Computer Engineering and Information Technology Department, Amirkabir University of Technology, Tehran, Iran.
- Kolla, B. P., Mansukhani, S., and Mansukhani, M. P. (2016). Consumer sleep tracking devices: a review of mechanisms, validity and utility. *Expert Review of Medical Devices*, 13(5):497–506.
- Kurihara, Y. and Watanabe, K. (2012). Sleep-stage decision algorithm by using heartbeat and body-movement signals. *IEEE Transactions on Systems, Man, and Cybernetics – Part A: Systems and Humans*, 42(6):1450–1459.
- Kwapisz, J. R., Weiss, G. M., and Moore, S. A. (2011). Activity recognition using cell phone accelerometers. *SIGKDD Explorations Newsletter*, 12(2):74–82.
- Långkvist, M., Karlsson, L., and Loutfi, A. (2012). Sleep stage classification using unsupervised feature learning. *Advances in Artificial Neural Systems*, 2012(3):1–9.
- O’Hare, E., Flanagan, D., Penzel, T., Garcia, C., Froberg, D., and Heneghan, C. (2015). A comparison of radio-frequency biomotion sensors and actigraphy versus polysomnography for the assessment of sleep in normal subjects. *Sleep and Breathing*, 19(1):91–98.
- Panagiotou, C., Samaras, I., Gialelis, J., Chondros, P., and Karadimas, D. (2015). A comparative study between SVM and fuzzy inference system for the automatic prediction of sleep stages and the assessment of sleep quality. In *Proc. 9th Int. Conf. on Pervasive Computing Technologies for Healthcare*, pages 293–296.
- Panossian, L. A. and Avidan, A. Y. (2009). Review of sleep disorders. *Medical Clinics of North America*, 93(2):407–425.
- Radha, M., Garcia-Molina, G., Poel, M., and Tononi, G. (2014). Comparison of feature and classifier algorithms for online automatic sleep staging based on a single EEG signal. In *Proc. 36th Annual Int. Conf. of the IEEE Engineering in Medicine and Biology Society*, pages 1876–1880.
- Rahman, T., Adams, A. T., Ravichandran, R. V., Zhang, M., Patel, S. N., Kientz, J. A., and Choudhury, T. (2015). Dopplesleep: A contactless unobtrusive sleep sensing system using short-range doppler radar. In *Proc. ACM Int. Joint Conf. on Pervasive and Ubiquitous Computing*, pages 39–50.
- Reimer, U. and Maier, E. (2016). An application framework for personalised and adaptive behavioural change support systems. In *Proc. 2nd Int. Conf. on Information and Communication Technologies for Ageing Well and e-Health (ICT4AWE)*.
- Reimer, U., Maier, E., Laurenzi, E., and Ulmer, T. (2017). Mobile stress recognition and relaxation support with SmartCoping: User adaptive interpretation of physiological stress parameters. In *Proc. Hawaii Int. Conf. on System Sciences (HICSS-50)*.
- Shi, J., Liu, X., Li, Y., Zhang, Q., Li, Y., and Ying, S. (2015). Multi-channel EEG-based sleep stage classification with joint collaborative representation and multiple kernel learning. *Journal of Neuroscience Methods*, 254:94–101.
- Sohm, M. (2016). *Erkennung von komplexen Aktivitäten anhand von tragbaren Sensoren*. Master thesis, University of Applied Sciences, Vorarlberg.
- Subramanian, S., Hesselbacher, S. E., Aguilar, R., and Surani, S. R. (2011). The NAMES assessment: a novel combined-modality screening tool for obstructive sleep apnea. *Sleep and Breathing*, 15(4):819–826.
- Tinguely, G., Landolt, H.-P., and Cajochen, C. (2014). Schlafgewohnheiten, Schlafqualität und Schlafmittelkonsum der Schweizer Bevölkerung: Ergebnisse aus einer neuen Umfrage bei einer repräsentativen Stichprobe. *Therapeutische Umschau*, 71(11):637–646.